



Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

Numerical Methods

EE1103

Indian Institute of Technology, Madras

February 9, 2021



References

Numerical
Methods

EE1103

Linux

Floats

Errors

μ , σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- Numerical Methods for Engineers, Chapra and Canale, 6th edition
- onlinegdb.com
- www.geeksforgeeks.org/c-programming-language/
- Linux cheat sheet



Intro to Linux

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- 1 Its an **Operating System**, like Windows or MacOS
- 2 It is **NOT** a programming language
- 3 It has different programs like Terminal, OpenOffice or Libre Office, Gnuplot,

There are many variants of Linux. The most popular ones being **Ubuntu, Redhat and Debian**

You can install **VirtualBox** and run multiple OS'es without having to reboot your computer.



The Linux Terminal

EE1103

Linux

Floats

Errors

μ , σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- The Terminal runs a **Shell** (default is likely to be **bash**)
- You can write bash-scripts that do interesting things
- You need to type commands at the prompt within a Terminal e.g.
 - **clear** will clear the terminal screen
 - **mkdir**, **cd**, **pwd**, and **popd/pushd** help you navigate within the folders/directories,
 - **ls** will list the files within a folder, **mv** will move them
 - **history** will give you the list of commands
 - you could also use e.g. **!3** or **!m** to recall commands from the history



The Linux Terminal

EE1103

Linux

Floats

Errors

μ , σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- You must learn to work with processes
 - `ps` (lists the processes, `fg` (foreground) or `bg` (background) allows you to control their execution
 - `Ctrl-C` : kill a running program `Ctrl-Z`: suspend a running program, restart with `fg` or `bg`
- Use the **man** pages to learn more about a command e.g.
 - **man ls**
- Redirection: overwrite (`>`), or append (`>>`) output from a command to a file e.g.
 - `ls -a > myfiles.txt`
- Use **more**, **less**, **tail** to look inside a file



Working with data

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- Generate your own data - try **ping** `www.iitm.ac.in`
- string editor `sed 's/old/new/g'` works on each line
 - the first **s** is to substitute
 - the last **g** stands for global
- `awk -F, '{print $7}' filename.txt`
 - uses the comma as a delimiter
 - prints the 7th column in the file
- Use `|` to pipe the output of a command to another command



Numerical Data Types

Numerical
Methods

EE1103

Linux

Floats

Errors

μ , σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- 1 Single float (float)
- 2 Double float (double)
- 3 Long Double float. (long double)

Called **floats** because the decimal point is moving.

C does not support **fixed-points**



Integer Representation

EE1103

Linux

Floats

Errors

μ , σ

Integration


Root Finding

Interpolation

ODE Solvers

Regression

- Binary
 - Powers of 2
- Decimal
 - Powers of 10
- Hexadecimal e.g. 0x1F
 - $A = 10$
 - $B = 11$
 - \vdots
 - $F = 15$





Numerical Data Types

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

15.875 can be converted to binary as

$$15 = 8 + 4 + 2 + 1 = \text{Binary}(1111)$$

$$.875 = 0.5 + 0.25 + 0.125 = 2^{-1} + 2^{-2} + 2^{-3}$$

So, the binary form of 15.875 is 1111.111 or represented as 01111.111



Space allocated

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- float: 4 bytes or 32 bits
- double: 8 bytes or 64 bits
- long double: 10 bytes or 80 bits

Learn to use the data type that is relevant.

Don't misuse long double, or double



32bit Float

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- 1 bit represent sign bit
- 8 bits represents exponent (bias is 127)
- 23 bits represents mantissa (recall characteristic and mantissa)
- 32 bit representation is 4 bytes

$15.875 =$ 01000001 01111110 00000000 00000000



32bit Float - Example

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

$$15.875 = 1111.111 = 1.111111 \times 2^3$$

- Number is +ve so sign bit is 0
- Mantissa (significand) is 111111 represented as 111111000000000000000000
- Exponent is 11 (or 3) represented as $\text{Binary}((127 + 3) = 130) = 10000010$
- 32 bit representation is 01000001011111110000000000000000 (4 bytes)

01000001 01111110 00000000 00000000



64bit Double

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- 1 bit represent sign bit
- 11 bits represent exponent (exponent bias is 1023)
- 52 bits represent mantissa.



128bit Long Double

Numerical
Methods

EE1103

Linux

Floats

Errors

μ , σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- 1 bit represent sign bit
- 15 bits represent exponent (exponent bias is 16383)
- 64 bits represent mantissa.

Check your conversions here

<https://www.h-schmidt.net/FloatConverter/IEEE754.html>



The Test

Numerical
Methods

EE1103

Linux

Floats

Errors

μ , σ

Integration

Root Finding

Interpolation

ODE Solvers


Regression

- Sample $\sin(x)$ for $x \in [0, 2\pi]$
- Define x as
 - float
 - double
 - long double
- Compare the values generated



Accurate and Precise

- Accurate: average value (mean) is correct
- Precise: spread (standard deviation) is small





Errors

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

$$\text{Error } \mathcal{E} = \frac{\text{Calculated Value} - \text{True Value}}{\text{True Value}}$$

- Round-off: using finite precision
- Truncation: numerical methods use approximations



Iterative Methods

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

$$e^{-x} = 1 - x + \frac{x^2}{2!} - \frac{x^3}{3!} + \dots$$

- Series must be truncated after some terms

$$\text{Normalized Error } \varepsilon = \frac{\text{approximate error}}{\text{approximation}} \times 100\%$$

$$\text{Relative Error } \varepsilon_a = \frac{\text{current approx} - \text{previous approx}}{\text{current approx}} \times 100\%$$



Taylor Series

Taylor Polynomial and remainder

$$f(x) = \sum_{n=0}^N \frac{1}{n!} f^{(n)}(c)(x-c)^n + \frac{1}{(N+1)!} f^{(N+1)}(\xi)(x-c)^{N+1}$$

Standard functions

$$e^x = 1 - x + \frac{x^2}{2!} - \frac{x^3}{3!} + \dots$$


$$\sin x = x - \frac{1}{3!}x^3 + \frac{1}{5!}x^5 + \dots$$

$$\ln(1+x) = x - \frac{1}{2}x^2 + \frac{1}{3}x^3 + \dots$$

$$(1+x)^p = 1 + px + \frac{p(p-1)}{2!} + \frac{p(p-1)(p-2)}{3!} + \dots$$



Probability Distributions



Probability Distribution: $Pr[a \leq X \leq b] = \int_a^b f_X dx$

Cumulative Distribution: $F_X = \int_{-\infty}^x f_X(u) du$



Normal Distribution

EE1103

Linux

Floats

Errors

μ, σ



Integration

Root Finding

Interpolation

ODE Solvers

Regression



$$f(x|\mu, \sigma^2) = \frac{1}{\sqrt{(2\pi\sigma^2)}} e^{-(x-\mu)^2/\sigma^2}$$

- Mean μ and variance σ^2
- Show that $\sigma^2 = \frac{1}{N} \sum_i (x_i - \langle x \rangle)^2 = \langle x^2 \rangle - \langle x \rangle^2$



Normal Distribution

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- Box-Mueller transform: generate a normal distribution from two uniform distributions U_1 and U_2

$$Z_0 = R \cos \theta = \sqrt{-2 \log U_1} \cos(2\pi U_2)$$

$$Z_1 = R \sin \theta = \sqrt{-2 \log U_1} \sin(2\pi U_2)$$

- Z_0 and Z_1 are independent variables with standard normal distributions
- https://en.wikipedia.org/wiki/Box-Muller_transform





Additive Noise

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- Generate a sine wave x as an array of points $x_{i=1\dots N}$
 - Inputs: amplitude V_{pp} , freq f , number of cycles M
 - Choose $3 < P < 10$ number of points per cycle such that $N = M \times P$
 - Add Gaussian noise e_i to create the noisy sine wave y_i
 - Choose $\mu = 0$, $\sigma = 0.1V_{pp}$
 - Use the Box-Muller transform to generate e_i
 - Estimate the error $\varepsilon_i = y_i - x_i$ for all N points and determine μ and σ of $\varepsilon_{i=1\dots N}$
 - Tabulate the statistics: what % of points lie within 2σ , 4σ , 6σ of the mean?
- Choose a random $\mu \neq 0$, $\sigma = 0.2V_{pp}$ and repeat the exercise
 - Does your choice of P and μ affect the distribution of e ?



Integral of a function

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

Find the value of $\int_a^b f(x)dx$ given a function $f(x) = 0$

- Rectangular bins
- Trapezoid rule
- Simpson's rule
- What are the errors?



Trapezoid Rule

EE1103

Linux

Floats

Errors

μ, σ


Integration

Root Finding

Interpolation

ODE Solvers

Regression




- Break the range $[x_L, x_U]$ into N equal bins
- Calculate

$$\begin{aligned} y_i &= \int_{x_i}^{x_{i+1}} f(x) dx \\ &= (\Delta x) \frac{f(x_{i+1}) - f(x_i)}{2} \end{aligned}$$

- Calculate $I_N = \sum_{i=1}^N y_i$
- Plot I_N versus N and look for **convergence**



Simpsons 1/3 Rule



- Break the range $[x_L, x_U]$ into N equal bins
- Approximate $f(x) \approx f_2(x)$

$$y_i = \int_{x_i}^{x_{i+1}} f_2(x) dx$$

- Calculate $I_N = \sum_{i=1}^N y_i$

- $f_2(x)$ is a second order *Lagrange polynomial*

$$y_i = (\Delta x) \frac{f(x_{i-1}) + 4f(x_i) + f(x_{i+1}))}{6}$$

- Plot I_N versus N and look for **convergence**



Finding the Root

Numerical
Methods

EE1103

Linux

Floats

Errors

μ , σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

Find the value of x that satisfies a given function $f(x) = 0$

- Simple fixed point iteration
- Bracketing
 - Bisection
 - False position
- Newton-Raphson
- Secant



Fixed Point Iteration

Numerical
Methods

EE1103

Linux

Floats

Errors

μ , σ

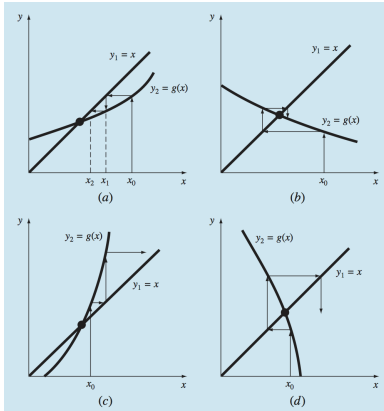
Integration

Root Finding

Interpolation

ODE Solvers

Regression



- Start at some value $x = x_0$
- Calculate $y_0 = f(x_0)$
- Set $x_1 = y_0$
- Iterate until
$$\varepsilon_i = \left| \frac{x_{i+1} - x_i}{x_{i+1}} \right| < \text{tolerance}$$

Try it for

- $f(x) = \sin(x) + x$
- $f(x) = e^{-x} - x$



Bracketing

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ


Integration

Root Finding

Interpolation

ODE Solvers

Regression



- Start with some values of x_L and x_U
- Calculate $y_L = f(x_L)$ and $y_U = f(x_U)$
- Check that $y_L \times y_U < 0$
- Estimate $x_i = 0.5(x_L + x_U)$
 - if $y_L f(x_i) < 0$ then $x_U = x_i$
 - else $x_L = x_i$
 - Define ε_i and iterate until $\varepsilon_i < \text{tolerance}$
- Find all the roots of the equation

$$f(x) = \sin 10x + \cos 3x \quad 0 < x < 5$$



False Position

Numerical
Methods

EE1103

Linux

Floats

Errors

μ , σ


Integration

Root Finding

Interpolation

ODE Solvers

Regression



- Start with some values of x_L and x_U
- Calculate $f(x_L)$ and $f(x_U)$
- Check that $f(x_L)f(x_U) < 0$
- Use the slope to estimate x_i

$$x_r = x_u - \frac{f(x_u)(x_l - x_u)}{f(x_l) - f(x_u)}$$

- Define ε_i and iterate until $\varepsilon_i < \text{tolerance}$



Linear Interpolation

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ


Integration

Root Finding

Interpolation

ODE Solvers

Regression



Fit $n + 1$ points to a
polynomial of order n

- Obtain the value $y = f(x)$ by interpolating between $y_0 = f(x_0)$ and $y_1 = f(x_1)$

$$y - y_0 = m(x - x_0) = \frac{(y_1 - y_0)}{(x_1 - x_0)}(x - x_0)$$



Polynomial Interpolation

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ


Integration

Root Finding

Interpolation

ODE Solvers

Regression



- Fit $n + 1$ points to a polynomial of order n
- Set of equations for points $(x_0, y_0) \dots (x_n, y_n)$



Vandermode Polynomial

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

$$P_n(x) = c_0x^n + c_1x^{n-1} + \dots + c_n$$

- System of equations to be solved

$$\begin{pmatrix} x_0^n & x_0^{n-1} & \dots & x_0 & 1 \\ x_1^n & x_1^{n-1} & \dots & x_1 & 1 \\ \vdots & \vdots & & \vdots & \vdots \\ x_n^n & x_n^{n-1} & \dots & x_n & 1 \end{pmatrix} \begin{pmatrix} c_0 \\ c_1 \\ \vdots \\ c_n \end{pmatrix} = \begin{pmatrix} y_0 \\ y_1 \\ \vdots \\ y_n \end{pmatrix}$$

- Easiest to do using matrix inversion methods



Lagrange Polynomial

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

$$\begin{aligned} P_n(x) = & \frac{(x - x_1)(x - x_2) \dots (x - x_n)y_0}{(x_0 - x_1)(x_0 - x_2) \dots (x_0 - x_n)} \\ & + \frac{(x - x_0)(x - x_2) \dots (x - x_n)y_1}{(x_1 - x_0)(x_1 - x_2) \dots (x_1 - x_n)} \\ & + \frac{(x - x_0)(x - x_1) \dots (x - x_n)y_2}{(x_2 - x_0)(x_2 - x_1) \dots (x_2 - x_n)} \\ & + \dots + \frac{(x - x_0)(x - x_1) \dots (x - x_{n-1})y_n}{(x_n - x_0)(x_n - x_1) \dots (x_n - x_{n-1})} \end{aligned}$$

Note: First term does not have $(x - x_0)$ in the numerator, and uses x_0 in the denominator. Similar for each subsequent term.



Newton Polynomial

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

$$P_n(x) = c_0 + c_1(x - x_0) + c_2(x - x_0)(x - x_1) \\ + \dots + c_n(x - x_0) \cdots (x - x_{n-1})$$

Substitutions at each given point (x_i, y_i) results in a lower diagonal matrix that is easy to invert

$$y_0 = c_0$$

$$y_1 = c_0 + c_1(x - x_0)$$

$$y_2 = c_0 + c_1(x - x_0) + c_2(x - x_0)(x - x_1)$$

$$\vdots$$

$$y_n = c_0 + c_1(x - x_0) + \dots + c_n(x - x_0) \cdots (x - x_{n-1})$$



Adaptive Runge Kutta

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- Step halving: Compare the solution of one large step with two half steps
- Embedded RK methods: Compare solutions from two different RK methods to detect sudden changes

$$y_{i+1} = y_i + \frac{1}{9} (2k_1 + 3k_2 + 4k_3) h$$

$$k_1 = f(t_i, y_i)$$

$$k_2 = f\left(t_i + \frac{1}{2}h, y_i + \frac{1}{2}k_1h\right)$$

$$k_3 = f\left(t_i + \frac{3}{4}h, y_i + \frac{3}{4}k_2h\right)$$

Local truncation error $E_{i+1} = \frac{1}{72} (-5k_1 + 6k_2 + 8k_3 - 9k_4) h$
where $k_4 = f(t_{i+1}, y_{i+1})$



Stiff Equations

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- Solutions have both a rapidly changing component and a slowly changing component
- Often the rapidly changing one's die as initial transients and the solution is dominated by the slow solution
- Example

$$\frac{dy}{dt} = -1000y + 3000 - 2000e^{-t}$$


- for $y(0) = 0$ we get the analytic solution

$$y(t) = 3 - 0.998e^{-1000t} - 2.002e^{-t}$$



Stiff Equations


- $y = 3 - 0.998e^{-1000t} - 2.002e^{-t}$
- Fast transient from $y = 0$ to $y = 1$ that occurs in $0 < t < 0.005$





Linear Regression

- Values of x are known without error
- Values of y are independent random variables with the same variance
- Values of y for a given x must be normally distributed





Linear Regression

Numerical
Methods

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- Given a data set $y(x)$
- Best fit is $y = a_0 + a_1x + e$
- Residual error $e = y - a_0 + a_1x$
- Minimize the sum of the squares of the residuals

$$S_r = \sum_{i=1}^n e_i^2 = \sum_{i=1}^n (y_{i,\text{data}} - y_{i,\text{fit}})^2 = \sum_{i=1}^n (y_i - a_0 - a_1x_i)^2$$



Linear Regression

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- Sum of the squares of the residuals

$$S_r = \sum_{i=1}^n e_i^2 = \sum_{i=1}^n (y_{i,\text{data}} - y_{i,\text{fit}})^2 = \sum_{i=1}^n (y_i - a_0 - a_1 x_i)^2$$

- Differentiate with respect to a_0 and a_1 and set to zero to minimize S

$$\begin{aligned} n a_0 + \left(\sum x_i \right) a_1 &= \sum y_i \\ \left(\sum x_i \right) a_0 + \left(\sum x_i^2 \right) a_1 &= \sum x_i y_i \end{aligned}$$

- Solve for a_0 and a_1



Goodness of Fit

EE1103

Linux

Floats

Errors

μ, σ

Integration

Root Finding

Interpolation

ODE Solvers

Regression

- How good is a straight line fit compared to just using the mean of the data?
- Normalize the residual error against the variance of the data

$$S_r = \sum_{i=1}^n e_i^2 = \sum_{i=1}^n (y_{i,\text{data}} - y_{i,\text{fit}})^2 = \sum_{i=1}^n (y_i - a_0 - a_1 x_i)^2$$

$$S_t = \sum_{i=1}^n (y_{i,\text{data}} - y_{\text{mean}})^2 = \sum_{i=1}^n y_i^2 - n y_{\text{mean}}^2$$

- Coefficient of determination


$$r^2 = \frac{S_t - S_r}{S_t}$$

- Correlation coefficient $0 < r < 1$
- **Unwise** to claim e.g. $r^2 = 0.93$ is better than $r^2 = 0.92$



Linearization


- Linearize before using regression





Poor Fits

- Errors must be normally distributed about the fit
- Many data sets will have **outliers**
 - more than 2σ away
 - decide whether to keep or remove them before you fit





Polynomial Regression

- Fit to $y_{\text{fit}} = a_0 + a_1x + a_2x^2 + e$
- Sum of the squares of the residuals

$$S_r = \sum_{i=1}^n e_i^2 = \sum_{i=1}^n (y_i - a_0 - a_1x_i - a_2x_i^2)^2$$

- Differentiate (paritials) with respect to a_0 , a_1 and a_2 and set them to zero to minimize S

$$na_0 + \left(\sum x_i\right) a_1 + \left(\sum x_i^2\right) a_2 = \sum y_i$$

$$\left(\sum x_i\right) a_0 + \left(\sum x_i^2\right) a_1 + \left(\sum x_i^3\right) a_2 = \sum x_i y_i$$

$$\left(\sum x_i^2\right) a_0 + \left(\sum x_i^3\right) a_1 + \left(\sum x_i^4\right) a_2 = \sum x_i^2 y_i$$

- Solve for a_0 , a_1 and a_2



Multiple Linear Regression

- Fit to $y_{\text{fit}} = a_0 + a_1x_1 + a_2x_2 + e$
- Sum of the squares of the residuals

$$S_r = \sum_{i=1}^n e_i^2 = \sum_{i=1}^n (y_i - a_0 - a_1x_{1i} - a_2x_{2i})^2$$

- Differentiate (partials) with respect to a_0 , a_1 and a_2 and set them to zero to minimize S

$$\begin{bmatrix} n & \sum x_{1i} & \sum x_{2i} \\ \sum x_{1i} & \sum x_{1i}^2 & \sum x_{1i}x_{2i} \\ \sum x_{2i} & \sum x_{1i}x_{2i} & \sum x_{2i}^2 \end{bmatrix} \begin{Bmatrix} a_0 \\ a_1 \\ a_2 \end{Bmatrix} = \begin{Bmatrix} \sum y_i \\ \sum x_{1i}y_{1i} \\ \sum x_{2i}y_{1i} \end{Bmatrix}$$

- Invert the coefficient matrix to solve for a_0 , a_1 and a_2



Standard Error

- Lose 2 degrees of freedom when we estimate a_0 and a_1
- Standard error of the estimate for a linear fit

$$s_{y/x} = \sqrt{\frac{S_r}{n-2}} \quad \text{for a fit to } n\text{-points}$$

- Polynomial fit, of order m , with $m+1$ DOF

$$s_{y/x} = \sqrt{\frac{S_r}{n-(m+1)}} \quad \text{for a fit to } n\text{-points}$$

- Also, if you fit to $y = a_0 + a_1x_1 + a_2x_2 + \cdots + a_mx_m + e$

$$s_{y/x} = \sqrt{\frac{S_r}{n-(m+1)}} \quad \text{for a fit to } n\text{-points}$$